



# DIRECTV - Enhanced Sports Replays

## AI Studio Final Presentation

Break Through Tech Los Angeles @UCLA  
December 4, 2023



# Introductions



# Meet Our Team!



**Rachel Cheung**  
UC Berkeley



**Joana Fang**  
UC Los Angeles



**Nhu Nguyen**  
CSU Long Beach



**Emi Viernes**  
CSU Northridge



# Our AI Studio TA and Challenge Advisors



**Kaifeng Pang**  
AI Studio TA



**Taylor Hosbach**  
Challenge Advisor



**Mia Calhoun Mummert**  
Challenge Advisor



# Presentation Agenda

1. AI Studio Project Overview
2. Data Understanding & Data Preparation
3. Modeling & Evaluation
4. Final Thoughts



# AI Studio Project Overview



# Identifying Football Players to Enhance Sports Replays





# Business Impact

- **70%** of DIRECTV customers watch at least **3 hours of sports each week**
- DIRECTV currently has Sports Mode
  - Enhanced Sports Replays would bring a more **personalized experience**
    - More engagement





# Sports Mode



Timeline

Team Comparison

Game Leaders

Enhanced Replays

BALL STATE  
**CARDINALS**  
10-2

**35**

4th | 05:36

**31**

TOLEDO  
**ROCKETS**  
4-7

## PASSING

J. Hall

J. Dart

276 YDS

20/32 COMP/ATT 23/35

248 YDS

2 TD 1  
2 INT 0

## RUSHING

T. Allgeier

V. Malepeai

21 ATT

111 YARDS 99

20 ATT

2 TD 1

## RECEIVING

T. Allgeier

G. Bryant Jr.

5 REC

35 YARDS 56

5 REC

0 TD 1



# Our Goal

1. Extract information from image data.
2. Finetune pre-trained models.
3. Evaluate model accuracy and select the best model.



## Resources We Leveraged



***matplot******lib***






# Data Understanding & Data Preparation



# Our Dataset from RoboFlow

The screenshot displays the RoboFlow website interface for the 'football-players-detection Image Dataset'. The top navigation bar includes links for Projects, Universe, Documentation, and Forum, along with a search bar and a 'Create Account' button. The main content area shows the dataset title, a 'Try Pre-Trained Model' button, and a 'Download Dataset' button. A 'VERSIONS' section lists three versions of the dataset, with the most recent version (2023-01-10 7:47pm) highlighted. Below this, 'POPULAR DOWNLOAD FORMATS' are listed, including YOLOv8, YOLOv5, YOLOv7, MT-YOLOv6, COCO JSON, YOLO Darknet, Pascal VOC XML, TFRecord, CreateML JSON, and Other Formats. The '255 Total Images' section shows a grid of image thumbnails with bounding boxes. The 'Dataset Split' section shows the distribution of images: TRAIN SET (204 Images), VALID SET (38 Images), and TEST SET (13 Images). The right sidebar features a 'SIMILAR PROJECTS' section and a 'More like' link to a related dataset.

- Given by DIRECTV
- Download data format: COCO JSON





# JSON Dataset

First half:

```
"categories": [{"id":1,"name":"ball"}, {"id":2,"name":"goalkeeper"},  
{"id":3,"name":"player"}, {"id":4,"name":"referee"}]  
, "images": [{"id":0,"license":1,"file_name":"121364_7_8_png.rf.d492  
bc91a93a35073d085c5110bc7c10.jpg","height":1280,"width":1280,"date  
_captured":"2023-08-07T22:52:26+00:00"}, ...
```

Second half:

```
"annotations": [{"id":0,"image_id":0,"category_id":2,"bbox":[0,465,  
10,49],"area":490,"segmentation":[],"iscrowd":0}, {"id":1,"image_id"  
:"id":0,"category_id":3,"bbox":[520,417,16.5,38],"area":627,"segmentat  
ion":[],"iscrowd":0}, {"id":2,"image_id":0,"category_id":3,"bbox":[  
623,412,13.5,47.5],"area":641.25,"segmentation":[],"iscrowd":0}, ..  
..
```



# JSON Dataset

First half:

```
"categories": [{"id":1,"name": "ball"}, {"id":2,"name": "goalkeeper"},  
{"id":3,"name": "player"}, {"id":4,"name": "referee"}]  
, "images": [{"id":0,"license":1,"file_name": "121364_7_8_png.rf.d492  
bc91a93a35073d085c5110bc7c10.jpg","height":1280,"width":1280,"date  
_captured": "2023-08-07T22:52:26+00:00"}, ...
```

Second half:

```
"annotations": [{"id":0,"image_id":0,"category_id":2,"bbox":[0,465,  
10,49],"area":490,"segmentation":[],"iscrowd":0}, {"id":1,"image_id  
":0,"category_id":3,"bbox":[520,417,16.5,38],"area":627,"segmentat  
ion":[],"iscrowd":0}, {"id":2,"image_id":0,"category_id":3,"bbox":[  
623,412,13.5,47.5],"area":641.25,"segmentation":[],"iscrowd":0}, ..  
..
```





# JSON Dataset

First half:

```
"categories": [{"id":1,"name":"ball"},{"id":2,"name":"goalkeeper"},  
{"id":3,"name":"player"},{"id":4,"name":"referee"}]  
, "images": [{"id":0,"license":1,"file_name":"121364_7_8_png.rf.d492  
bc91a93a35073d085c5110bc7c10.jpg","height":1280,"width":1280,"date  
_captured":"2023-08-07T22:52:26+00:00"}, ...
```

Second half:

```
"annotations": [{"id":0,"image_id":0,"category_id":2,"bbox":[0,465,  
10,49],"area":490,"segmentation":[],"iscrowd":0},{"id":1,"image_id  
":0,"category_id":3,"bbox":[520,417,16.5,38],"area":627,"segmentat  
ion":[],"iscrowd":0},{"id":2,"image_id":0,"category_id":3,"bbox":[  
623,412,13.5,47.5],"area":641.25,"segmentation":[],"iscrowd":0},..  
..
```



# JSON Dataset

First half:

```
"categories": [{"id":1,"name": "ball"}, {"id":2,"name": "goalkeeper"},  
{"id":3,"name": "player"}, {"id":4,"name": "referee"}]  
, "images": [{"id":0,"license":1,"file_name": "121364_7_8_png.rf.d492  
bc91a93a35073d085c5110bc7c10.jpg","height":1280,"width":1280,"date  
_captured": "2023-08-07T22:52:26+00:00"}, ...
```

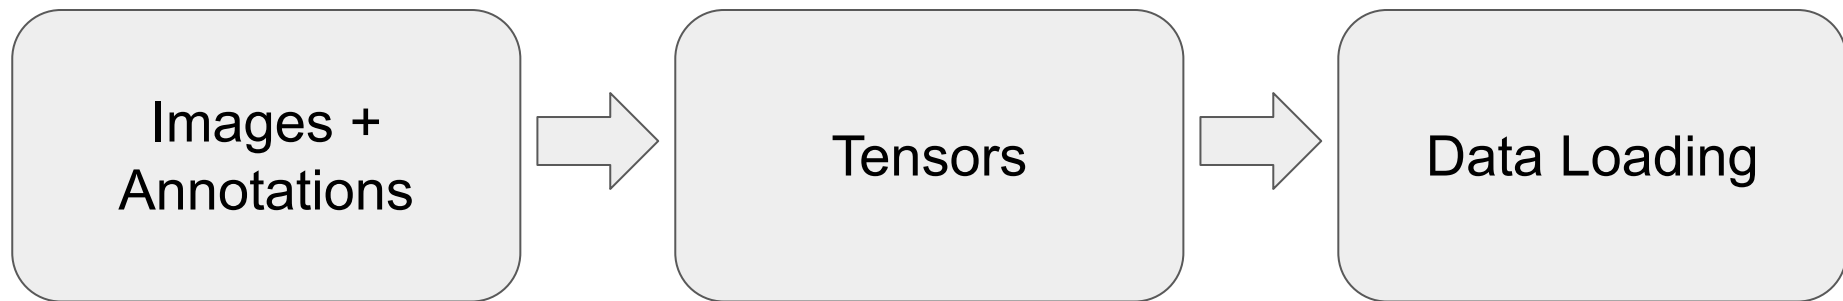
Second half:

```
"annotations": [{"id":0,"image_id":0,"category_id":2,"bbox": [0,465,  
10,49],"area":490,"segmentation": [],"iscrowd":0}, {"id":1,"image_id  
":0,"category_id":3,"bbox": [520,417,16.5,38],"area":627,"segmentat  
ion": [],"iscrowd":0}, {"id":2,"image_id":0,"category_id":3,"bbox": [  
623,412,13.5,47.5],"area":641.25,"segmentation": [],"iscrowd":0}, ..  
..
```



# Football Dataset Class and Data Loading

- Implemented functions to convert given images and their annotations to tensors
  - Tensors: multidimensional numerical arrays
- Loaded the data





# Modeling & Evaluation




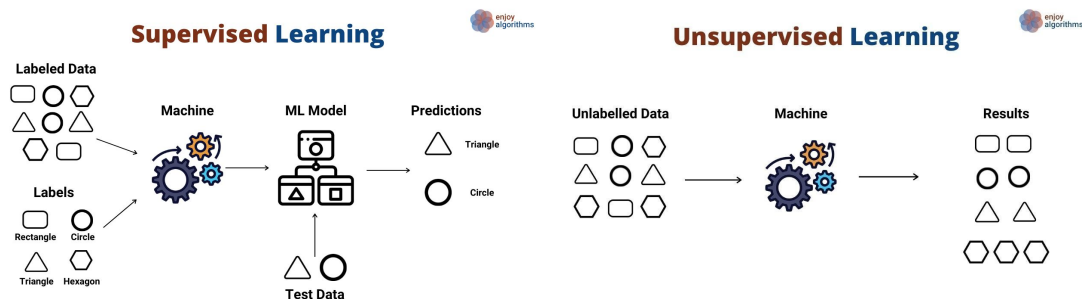
# Model Approach

Two types of Machine Learning processes: Supervised Learning vs Unsupervised

- Ours is supervised learning
  - Utilize the labeled data (JSON annotations) to train model and have the model output labeled data
  - Classification problem: our model should classify each object in the image to its respective category/class (player, goalkeeper, ball, and referee).
  - Utilize classification algorithms/methods

Classification vs Object Detection Algorithms (classification isn't enough)

- Classification algorithms: Take the whole image to classify object
- Object detection algorithms: Utilize bounding boxes (boxes around objects) within the image, more specific





# Models: R-CNN vs YOLO

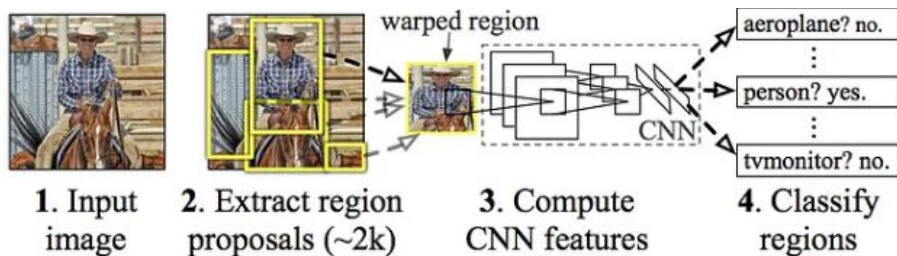
Existing popular object detection models we are using:

- **R-CNN** (Region-Based Convolutional Neural Network)
- **YOLO** (You Only Look Once)
  - Both utilize **CNN** (Convolutional Neural Network), more about this in the next slide

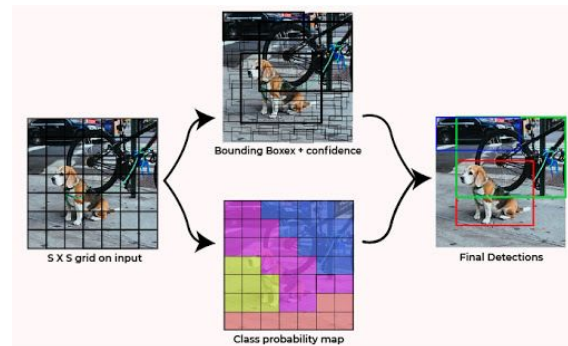
R-CNN vs YOLO

- R-CNN is a two-stage model: Propose regions with objects and then classify the objects in those regions
  - More steps = slower
- YOLO is a single-stage model: Predict objects and their bounding boxes at one look (You Only Look Once)
  - Less steps = faster

## R-CNN




## YOLO





# CNN (Convolutional Neural Network)

Neural Network



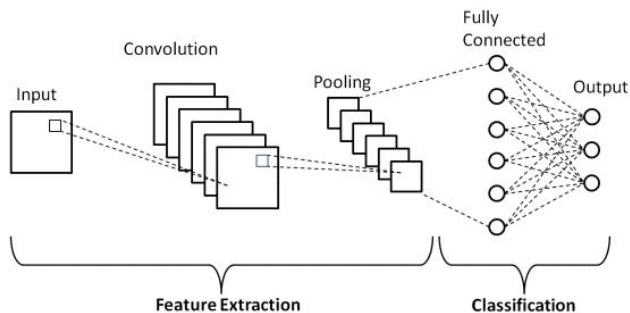
Input layer: Capture pixel data/features from the image like color of pixel

Hidden Layers: Recognize patterns in the pictures (like edges, patterns, complex shapes, etc, with each subsequent layer getting more complex) and deciding which features of the image are more important (weights)

Output Layer: The final outputs (the final predicted class and the score of how accurate it is) like 80% sure it is bird and 20% cat.



CNN



Conv layers/hidden: reduce images into a matrix of pixels that is easier to process (conv feature map) and define the weights

Pooling layers/hidden: simplifying data by keeping only the most important features

FC (fully connected)/hidden: connect all the features in the pooled feature maps

Softmax/Output Layer: convert raw scores to probability that we can use (in the format of 0.6, 0.7, etc.)

(conv feature map)

4	3	4
2	4	3
2	3	4



# Intro to R-CNN Family Models

Three generations:


- R-CNN
- Fast R-CNN
- **Faster R-CNN (the model we chose, more on the next slide)**

Basic three steps for all R-CNN models:

1. Create **region proposals** (2000) using a selective search algorithm: the regions in image that might contain object (bounding boxes)
2. Extract **CNN features** (patterns in the image like colors, lighting, etc.) from region proposals
3. Classify objects using features


The problem with R-CNN and Fast R-CNN: **they are slow** because they utilize **selective search algorithm** (an external algorithm) to generate region proposals

## RCNN



Region proposals, then CNN and then classification

## Fast RCNN



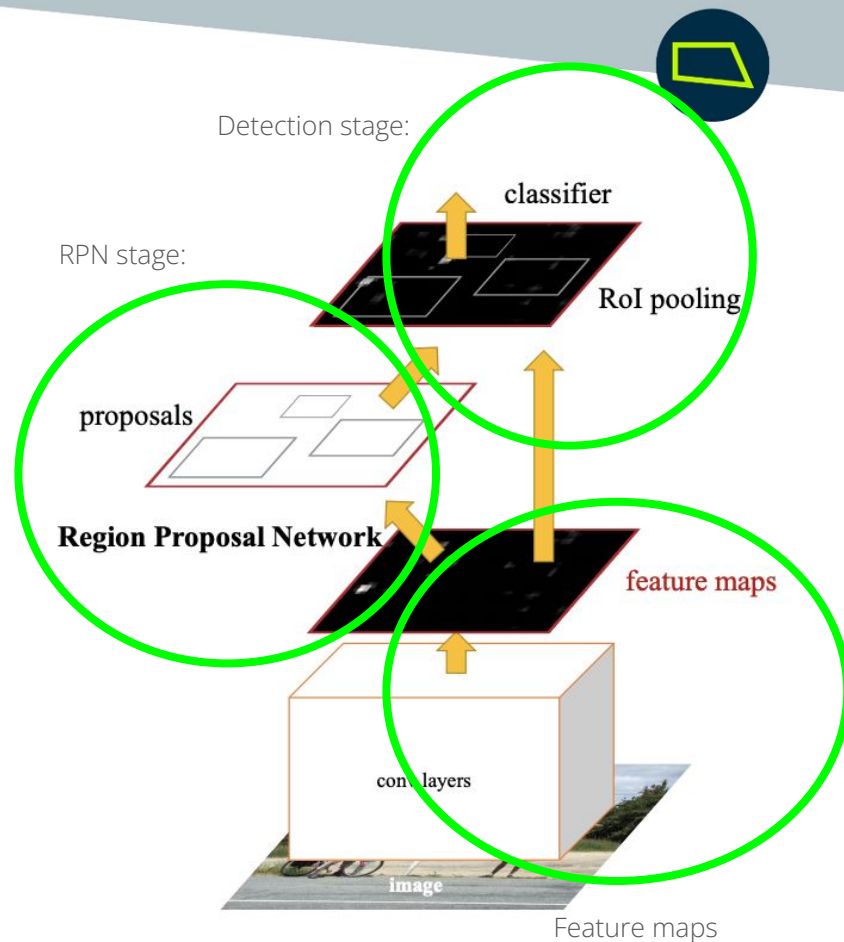
CNN (ConvNet + ROI/pool) with region proposals (step 2 and 3 combined) and then classification



# Faster R-CNN Model

Why Faster R-CNN? The fastest out of all RCNN models, allowing for real-time object detection.

- Region proposal network (RPN): Propose regions directly from the feature maps (building on the integration of CNN and region proposals in Fast R-CNN) and replacing selective search algorithm
- Detects objects with the proposed regions through ROI pooling and classifier
  - ROI pooling: makes sure all regions are reduced into fixed-size and only keep important features, reducing computational costs (better than Fast R-CNN by doing ROI pooling once)
  - Classifier: identify object label/class (player, referee, etc.) and refine coordinates of bounding box
- Shared CNN Backbone: Allow for model to utilize the same feature extractions in both the RPN stage and the detection stage (we can see the feature maps being used for both stages).

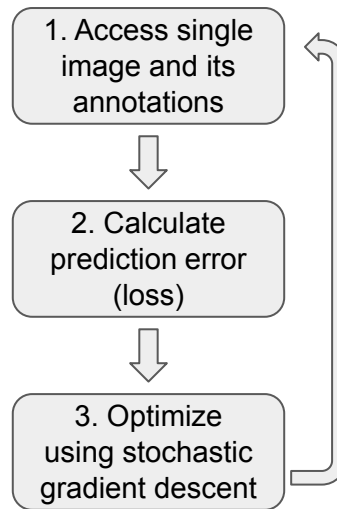




# Training Faster-RCNN Model

Steps:

1. Train all 600+ images in our training dataset with model by loading images and annotations
2. Have the model predict where the bounding boxes are and calculate the error between the prediction output and the provided target value (our annotations) to get loss
3. With our loss, optimize the model using stochastic gradient descent for each image to adjust the parameters of the model (the weights of each feature) and minimize loss (the learning step!)
4. And loop! In our case, we trained 600+ images 10 times (or epochs = passings of the entire dataset) for the best results.



```
Iteration: 1/10, Average iteration loss: 1.1697361469268799
Iteration: 2/10, Average iteration loss: 0.8446316123008728
Iteration: 3/10, Average iteration loss: 0.6970969438552856
Iteration: 4/10, Average iteration loss: 0.5963253378868103
Iteration: 5/10, Average iteration loss: 0.5403541326522827
Iteration: 6/10, Average iteration loss: 0.4978560507297516
Iteration: 7/10, Average iteration loss: 0.4836512506008148
Iteration: 8/10, Average iteration loss: 0.42549413442611694
Iteration: 9/10, Average iteration loss: 0.46847689151763916
Iteration: 10/10, Average iteration loss: 0.4042800962924957
```




# Output of Faster RCNN Model

Snippet of output predictions  
for 13 test images

```
[{'boxes': tensor([[ 325.3285,  
470.2320,  340.2840,  519.5388]]),  
device='cuda:0'), 'labels':  
tensor([2, 2, 2, 2]),  
device='cuda:0'), 'scores':  
tensor([0.9995, 0.9992, 0.9991,  
0.9991]), device='cuda:0')}]
```

Draw function




Actual labels

Predictions



# Evaluate Faster R-CNN Model: Intersection over Union (IoU metrics)

- Calculated accuracy scores with intersection over union
  - Compared predicted boxes with actual boxes
  - Higher IoU score = better accuracy



```
{'iou': tensor(0.0471, device='cuda:0')}
```

```
{'iou': tensor(0.0447, device='cuda:0')}
```

```
{'iou': tensor(0.0331, device='cuda:0')}
```




# NMS

- Problem: Faster R-CNN can produce multiple bounding boxes for one object
- Solution: **NMS** (Non-maximum Suppression)
  - Reduces the number of output bounding boxes using IOU by iteratively removing boxes with low scores and IOUs that are bigger than the IOU threshold we want (small IOU means keeping only the most precise boxes)


```
indices = torchvision.ops.batched_nms(bboxes=pred[0]['boxes'], scores=pred[0]['scores'], idxs = pred[0]['labels'],  
iou_threshold = 0.2) *WE WANT LOW IOU THRESHOLD (but at some point, lowering the iou threshold will not increase score  
because there are no overlapping boxes left)*
```

Before



```
{'iou': tensor(0.0436, device='cuda:0')}
```

After



```
{'iou': tensor(0.0461, device='cuda:0')}
```




# More NMS

Actual Labels



```
{'iou': tensor(0.0591, device='cuda:0')}
```

NMS Predictions



```
{'iou': tensor(0.0461, device='cuda:0')}
```

- Problem with our implementation of NMS: potential overlapping between two classes (minimize class bounding boxes), strange IOU scores for actual labels
- Problem with this model training: can't detect ball at all



What About YOLO?






# YOLO: Real-Time Object Detection

Goal: Create a **YOLOv8** model with **Darknet** framework and finetune it to our specific 4 classes

- Why YOLO? YOLO is extremely **fast** and **accurate**
  - Speed vs accuracy adjustments are made via **changing** model **size**
  - **NO RETRAINING!!!!**



SAY

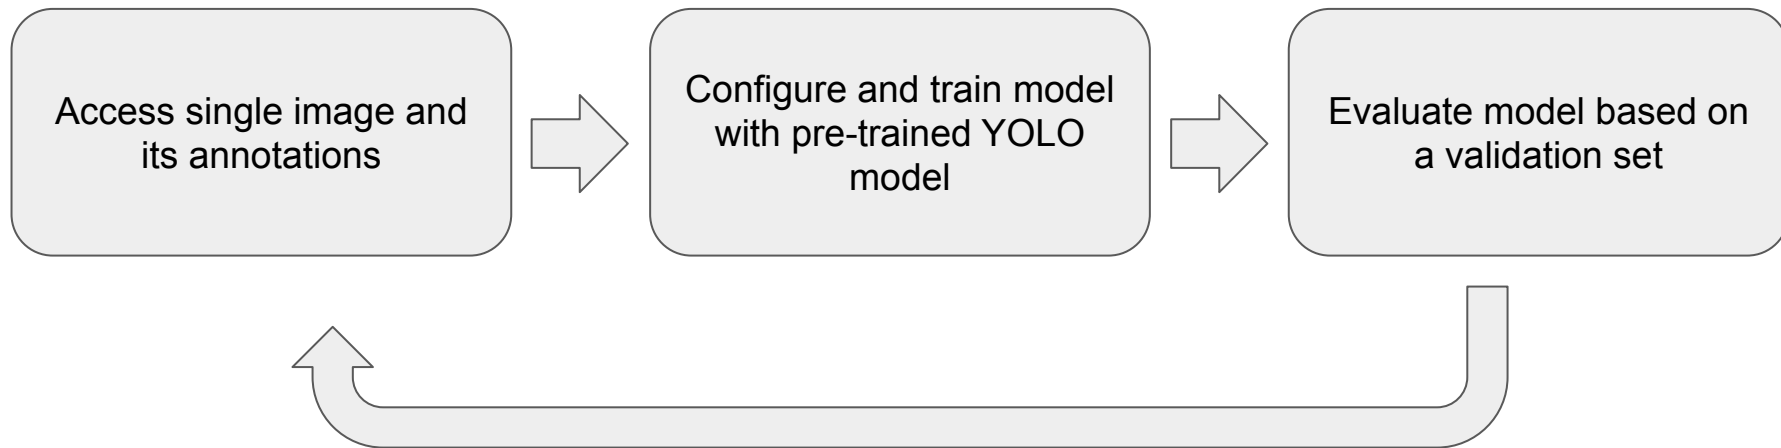
YOLO

AGAIN





# Training YOLO Model

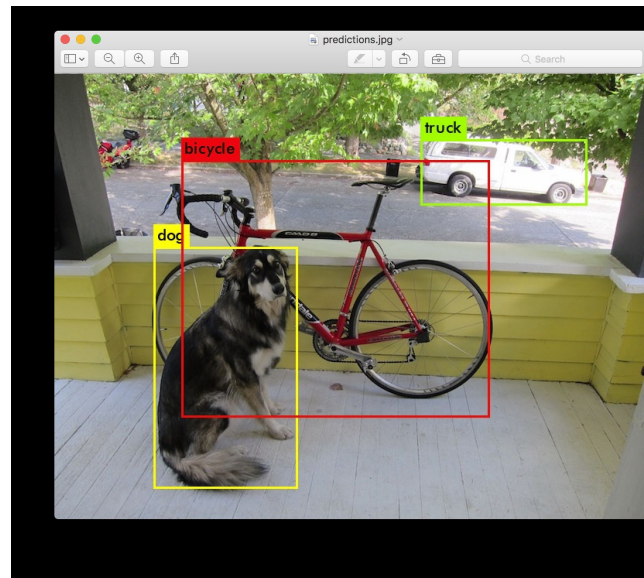




# Darknet framework used alongside YOLO

Utilizing **darknet** over a **traditional** approach

- Why darknet? Utilizes **different approach**
  - **Standard**: Repurpose **classifiers** or **localizers** by applying the model at multiple **locations** and **scales** of images, **region scoring** for detection
  - **Standard (R-CNN)**: Makes predictions that require thousands of images of neural networks for 1 image R
  - **Darknet**: makes it **extremely fast**, more than 1000x faster than R-CNN and 100x faster than Fast CNN



# Evaluation of YOLO Model

## PT.1-Understanding Format

- YOLO Formatting- **NO JSON allowed**
  - Converted into one row per object in **class x\_center y\_center width height format**
- Very specific pathing needed when importing our data set


### SAMPLE SET



```
# Classes (80 COCO classes)
names:
 0: person
 1: bicycle
 2: car
...
77: teddy bear
78: hair drier
79: toothbrush
```



```
1 0.384375 0.277734375 0.00625 0.01015625
3 0.44453125 0.2294921875 0.0109375 0.043359375
4 0.2396484375 0.2224609375 0.012109375 0.043359375
```





# Evaluation of YOLO Model PT.2- YOLO on a single Image

- Exploratory Test
  - Checked how information would be outputted
  - First look at the accuracies

```
Loading weights from yolov3.weights...Done!  
/content/data/test/40cd38_7_6_png.rf.dc7957fcce915b  
sports ball: 54%  
person: 100%  
person: 99%  
person: 98%  
person: 97%  
person: 97%  
person: 97%  
person: 95%  
person: 95%  
person: 94%  
person: 87%  
person: 86%  
person: 79%  
person: 77%  
person: 76%  
person: 75%  
person: 69%  
person: 69%  
person: 67%
```



## Evaluation of YOLO Model PT.3- Yolo Darknet Player Detection

- Model detection
  - Before anything the cfg coco.data must be specified to fit your pathing given names and number of classes.
  - The yolo.cfg model is modified for training
  - Then you are ready to train

```
./darknet detector train cfg/coco.data  
cfg/yolov3.cfg darknet53.conv.74
```

```
Region 94 Avg IOU: 0.244487, Class: 0.641122, Obj: 0.508777, No Obj: 0.518971, .5R:
```

```
Region 106 Avg IOU: 0.206882, Class: 0.393139, Obj: 0.472929, No Obj: 0.416488, .5R: 0.063
```

```
[net]  
# Testing  
# batch=1  
# subdivisions=1  
# Training  
batch=64  
subdivisions=8  
....
```

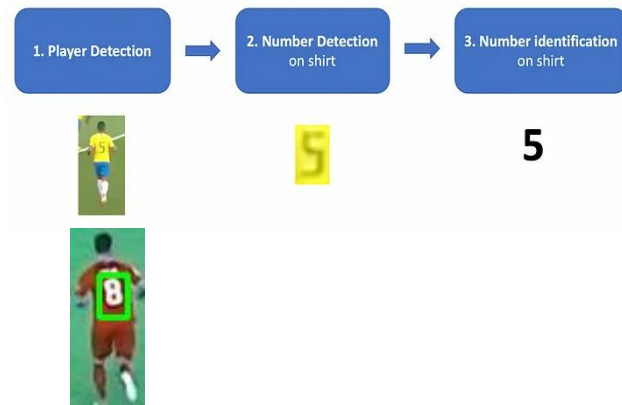
```
classes= 80  
train = <path-to-coco>/  
valid = <path-to-coco>/  
names = data/coco.names  
backup = backup
```



## Evaluation of YOLO Model PT.4- Yolo Darknet Detection Evaluation

- 3 main components to be created for complete player tracking in something like a video or frame
  1. Player detection
  2. Number detection
  3. Number identification
- BBox regression loss on training = large component of prediction like prev models
  - Comparing **Epoch** count to **Loss**
- Confusion Matrix Heat Maps
  - Later Allowed for easy prediction visualization of numbers to results

Fig. 3 Main Components of the YOLO notebook





# Model Comparison

Model Name	Description	Results	Pros	Cons
Faster-RCNN	Pre-trained deep network model that can predict object boxes and scores	Slow training but accurate predictions (with NMS) in relative to actual IOU	Identified the main objects correctly, fastest out of all R-CNN models	Multiple bounding boxes for one object, can't identify ball, slow when doing multiple epoch/iterations for large datasets
YOLOv8	Computer vision model for object detection and classification	Higher initial loss with IOU sadly below threshold	Worked fast, weights are able to help identify all the pretrained class names in images	Constant pathing and formatting issues, accuracy drops with object size



Final Thoughts





# What We Learned

- New tools, concepts, technologies, and frameworks
  - Google Colab, COCO JSON file + API, tensors, PyTorch, Darknet, IoU, NMS
- Object detection using pre-trained models (Faster-RCNN, YOLO)
  - First experience with **computer vision!**
- How to approach projects with fewer specifications



# Takeaways:

- Importance of organization and coordination
- Understand the tools and their limitations
- Utilize online resources
- Don't be afraid to ask for help



# Potential Next Steps

Models:

- Finetune R-CNN NMS more (there are still some overlaps, particularly for different classes)
- Implement R-CNN IOU better: strange IOU scores for actual labels

More complicated usage:

- Track players based on jersey number
- Detect objects on videos
- Specific replay angles based on player



Questions?